



Acoustic Modeling and Neural Network

Instructor: Tom Ko



The emission probabilities

- ▶ Recall that there is a $b_j(x_t)$ in the Viterbi algorithm and forward-backward algorithm
- ▶ $b_j(x_t) = p(x_t | s_j)$
- ▶ This represents the relationship between the audio signal and the phonetic units and is modeled by the acoustic model.



Start from the most simplest case


- ▶ Assume that there are only two sound units to classify: AO(啊) and AA(呀)
- ▶ Let x be one of the testing samples, we want to compute $P(\text{AO}|x)$ and $P(\text{AA}|x)$

$$P(\omega_j|x) = \frac{p(x|\omega_j)P(\omega_j)}{p(x)},$$

- ▶ Thus we need $P(x|\text{AO})P(\text{AO})$ and $P(x|\text{AA})P(\text{AA})$
- ▶ For the likelihood terms, we use Gaussian models:


$$p(x|\omega_1) = \frac{1}{\sqrt{2\pi}\sigma_1} \exp \left[-\frac{1}{2} \left(\frac{x - \mu_1}{\sigma_1} \right)^2 \right]$$

Classification with GMMs



- ▶ Obtain $P(\text{AO})$ and $P(\text{AA})$ by simple counting
- ▶ Build a Gaussian model for each class: $N(\mu_1, \sigma_1^2)$ and $N(\mu_2, \sigma_2^2)$
- ▶ Compute $P(x|\text{AO})$ and $P(x|\text{AA})$
- ▶ Make decision depending on $P(\text{AO}|x)$ and $P(\text{AA}|x)$

Classification with a neural network



- ▶ The output is directly the posterior $p(s_j | x_t)$
- ▶ .

Generative vs. discriminative

- ▶ Let's say you have input data x and you want to classify the data into labels y .
- ▶ A generative model learns the **joint** probability distribution $p(x,y)$
- ▶ A discriminative model learns the **conditional** probability distribution $p(y|x)$ – which you should read as *"the probability of y given x "*.
- ▶ Suppose you have the following data in the form (x,y) :
- ▶ $(1,0), (1,0), (2,0), (2, 1), (1,2), (1,2)$

- ▶ $p(x,y)$ is

	$y=0$	$y=1$	$y=2$
$x=1$	$2/6$	0	$2/6$
$x=2$	$1/6$	$1/6$	0

- ▶ $p(y|x)$ is

	$y=0$	$y=1$	$y=2$
$x=1$	$1/2$	0	$1/2$
$x=2$	$1/2$	$1/2$	0



Generative vs. discriminative

- ▶ GMM is a generative model
 - Each GMM is built for each class separately
 - It expresses how to generate the features if we knew it was of a particular class.
- ▶ NN is a discriminative model
 - It will learn to assign high weight to features that directly improve its ability to discriminate between possible classes
- ▶ The overall observation is that discriminative models generally outperform generative models in classification tasks.



Generative vs. discriminative

- ▶ In Bayesian learning, the posterior probability is transformed into a likelihood and a prior. They are learnt separately.
 - Thus, it is a generative classifier.
- ▶ How about a discriminative classifier which learn and compute the posterior probability directly?
 - Logistic regression is a discriminative classifier
 - A neural network can be viewed as a series of logistic regression classifiers stacked on top of each other



Logistic regression

- ▶ Consider a single input observation x , which is a vector of features $[x_1, \dots, x_n]$
- ▶ Logistic regression solves this task by learning, from a training set, a vector of weights and a bias term.

$$z = \left(\sum_{i=1}^n w_i x_i \right) + b$$

- ▶ The resulting single number z expresses the weighted sum of the evidence for the class.

$$z = w \cdot x + b$$

The sigmoid classifier

- ▶ Let us first assume a two-class classification problem.
- ▶ We'll pass z through the sigmoid function $\sigma(z)$
- ▶ The sigmoid function (named because it looks like an s) is also called the logistic function





Figure 5.1 The sigmoid function $y = \frac{1}{1+e^{-z}}$ takes a real value and maps it to the range $[0, 1]$. Because it is nearly linear around 0 but has a sharp slope toward the ends, it tends to squash outlier values toward 0 or 1.

The sigmoid classifier





The sigmoid classifier

- ▶ The sigmoid has a number of advantages:
 - It takes a real-valued number and maps it into the range $[0;1]$, which is just what we want for a probability.
 - Because it is nearly linear around 0 but has a sharp slope toward the ends, it tends to squash outlier values toward 0 or 1.
 - It's differentiable
- ▶ If $\sigma(z)$ represents the probability $P(w_1|x)$, then $1-\sigma(z)$ represents the probability $P(w_2|x)$
- ▶ For a test instance x , we say w_1 if the probability $\sigma(z)$ is more than 0.5, and no otherwise.
- ▶ We call 0.5 the decision boundary



Loss function

- ▶ We need a loss function that expresses, for an observation x , how close the classifier output ($\hat{y} = \sigma(z)$) is to the correct output (y , which is 0 or 1). We'll call this:

$$L(\hat{y}, y) = \text{How much } \hat{y} \text{ differs from the true } y$$



Loss function

- ▶ The probability of class w_1 ($y=1$, the positive class) computed by the model given an observation x is \hat{y}
- ▶ Similarly, the probability of class w_2 ($y=0$, the negative class) computed by the model is $(1 - \hat{y})$
- ▶ Thus, we can express $p(y|x)$ as

$$p(y|x) = \hat{y}^y (1 - \hat{y})^{1-y}$$

$$\begin{aligned}\log p(y|x) &= \log [\hat{y}^y (1 - \hat{y})^{1-y}] \\ &= y \log \hat{y} + (1 - y) \log (1 - \hat{y})\end{aligned}$$

- ▶ This is the term that we want to maximize.
- ▶ By flipping the sign, this is just happened to be the cross-entropy loss.



Cross-entropy loss function

- ▶ The cross-entropy loss is defined to be

$$L_{CE}(\hat{y}, y) = -[y \log \hat{y} + (1 - y) \log(1 - \hat{y})]$$

- ▶ By substituting $\hat{y} = \sigma(w \cdot x + b)$


$$L_{CE}(w, b) = -[y \log \sigma(w \cdot x + b) + (1 - y) \log(1 - \sigma(w \cdot x + b))]$$

- ▶ The total loss of the whole training set:

$$\sum_{i=1}^m L_{CE}(\hat{y}^{(i)}, y^{(i)})$$

Gradient descent


$$\hat{\theta} = \operatorname{argmin}_{\theta} \frac{1}{m} \sum_{i=1}^m L_{CE}(y^{(i)}, x^{(i)}; \theta)$$



$$\theta_{t+1} = \theta_t - \eta \nabla L(f(x; \theta), y)$$


How a sigmoid classifier classifies

- ▶ $w \cdot x = |w| |x| \cos\theta$
- ▶ We assume the bias term is 0 in this case, it will shift the decision boundary for non zero values.





How it changes in training




○ positive
+ negative



How it changes in training


- ▶ w will draw close to positive samples



○ positive
+ negative




How it changes in training



○ positive
+ negative

How it changes in training

- ▶ w will draw away from negative samples





Softmax Classifier

- ▶ Softmax classifier is a generalization of sigmoid classifier when the number of classes > 2 .
- ▶ It is also called multinomial logistic regression
- ▶ For a vector z of dimensionality k , the softmax is defined as:

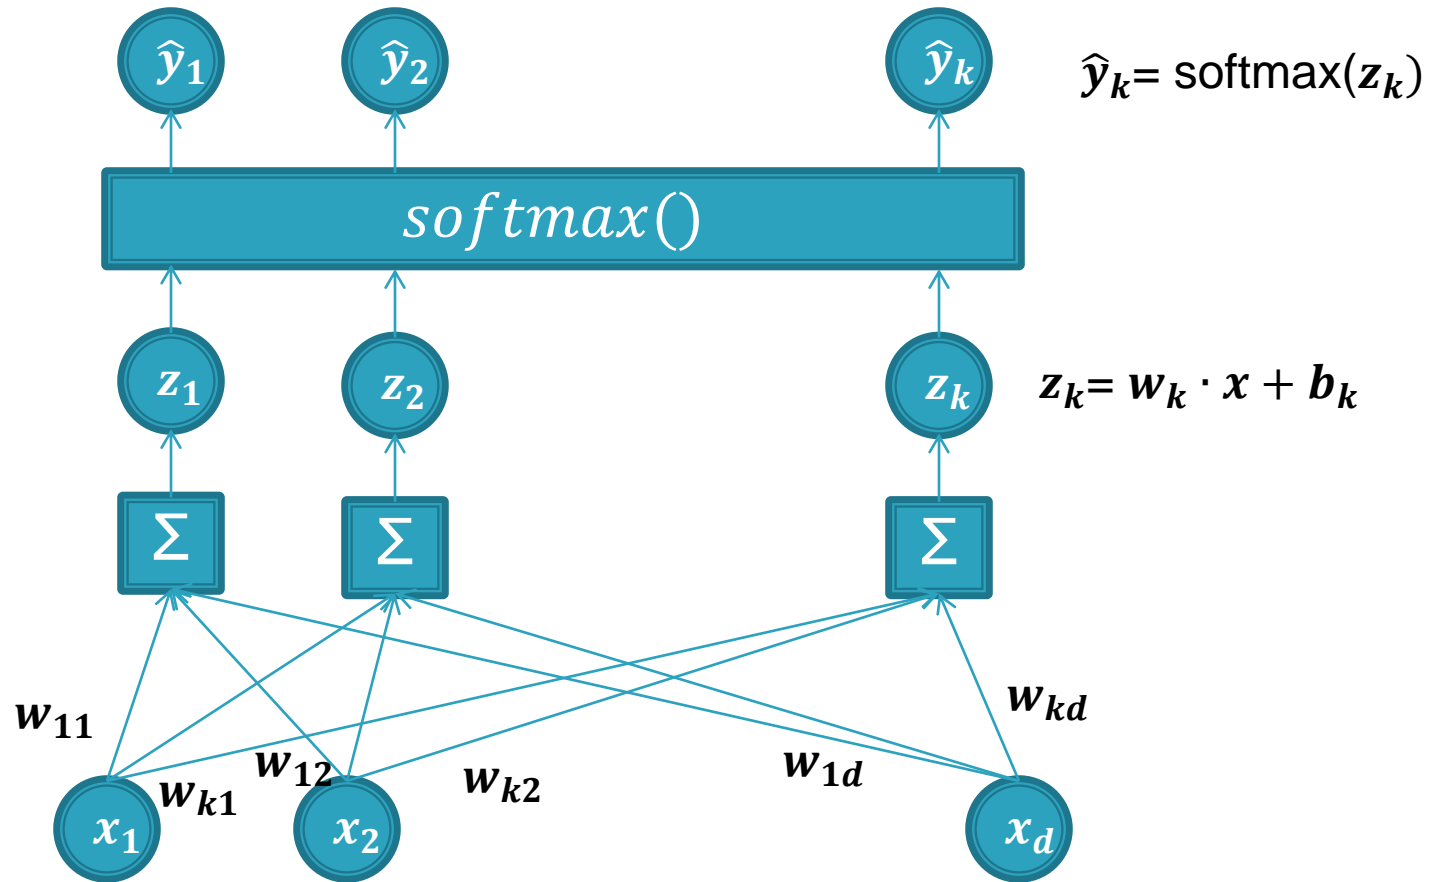
$$\text{softmax}(z_i) = \frac{e^{z_i}}{\sum_{j=1}^k e^{z_j}} \quad 1 \leq i \leq k$$

- ▶ The probability of class c among K classes $P(w_c|x) =$

$$\frac{e^{w_c \cdot x + b_c}}{\sum_{j=1}^k e^{w_j \cdot x + b_j}}$$


- ▶ The outputs of a softmax sum up to 1

The softmax classifier



How a softmax classifier classifies

► $w \cdot x = |w| |x| \cos\theta$





Derivative of softmax

- ▶ $\frac{\partial \hat{y}_k}{\partial z_i} = \frac{\partial \frac{e^{z_k}}{\sum e^{z_j}}}{\partial z_i} = \frac{\sum e^{z_j} \frac{\partial e^{z_k}}{\partial z_i} - e^{z_k} \frac{\partial \sum e^{z_j}}{\partial z_i}}{[\sum e^{z_j}]^2}$
- ▶ If $i = k$, the above term becomes
- ▶ $\frac{\sum e^{z_j} e^{z_k} - e^{z_k} e^{z_k}}{[\sum e^{z_j}]^2} = \frac{e^{z_k} (\sum e^{z_j} - e^{z_k})}{[\sum e^{z_j}]^2}$
- ▶ $= \frac{e^{z_k}}{\sum e^{z_j}} \frac{\sum e^{z_j} - e^{z_k}}{\sum e^{z_j}} = \frac{e^{z_k}}{\sum e^{z_j}} \left(1 - \frac{e^{z_k}}{\sum e^{z_j}}\right) = \hat{y}_k (1 - \hat{y}_k)$
- ▶ If $i \neq k$, the term becomes
- ▶ $\frac{-e^{z_k} e^{z_i}}{[\sum e^{z_j}]^2} = -\frac{e^{z_k}}{\sum e^{z_j}} \frac{e^{z_i}}{\sum e^{z_j}} = -\hat{y}_k \hat{y}_i$

Cross-entropy loss for multi class

- ▶ The cross-entropy loss is defined to be

$$L_{CE}(\hat{y}, y) = - \sum_{k=1}^K y_k \log \hat{y}_k$$

- where

$$\hat{y}_k = \frac{e^{w_k \cdot x + b_k}}{\sum_{j=1}^K e^{w_j \cdot x + b_j}}$$

- ▶ Thus,

$$L_{CE}(\hat{y}, y) = - \sum_{k=1}^K y_k \log \frac{e^{w_k \cdot x + b_k}}{\sum_{j=1}^K e^{w_j \cdot x + b_j}}$$

$$y_k = \begin{cases} 1 & \text{if } x \text{ is sample of class } k \\ 0 & \text{otherwise} \end{cases}$$



Learning in softmax classifier

- ▶ $\frac{\partial y_k}{\partial z_i} = \frac{\partial \frac{e^{z_k}}{\sum e^{z_j}}}{\partial z_i} = \frac{\sum e^{z_j} \frac{\partial e^{z_k}}{\partial z_i} - e^{z_k} \frac{\partial \sum e^{z_j}}{\partial z_i}}{[\sum e^{z_j}]^2}$
- ▶ Consider a class k sample, the CE loss becomes $-\log \hat{y}_k$
- ▶ $\frac{\partial L_{CE}}{\partial w_k} = - \frac{\partial \log \hat{y}_k}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial z_k} \frac{\partial z_k}{\partial w_k} = - \frac{1}{\hat{y}_k} \hat{y}_k (1 - \hat{y}_k) x = (\hat{y}_k - 1) x$
- ▶ Consider a class i (i \neq k) sample, the CE loss becomes $-\log \hat{y}_i$
- ▶ $\frac{\partial L_{CE}}{\partial w_k} = - \frac{\partial \log \hat{y}_i}{\partial \hat{y}_i} \frac{\partial \hat{y}_i}{\partial z_k} \frac{\partial z_k}{\partial w_k} = - \frac{1}{\hat{y}_i} (- \hat{y}_k \hat{y}_i) x = \hat{y}_k x$
- ▶ For a set of training data, the gradients for w_k are accumulated



Update the parameters

- ▶ The parameters will be updated according to

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{d}{d\mathbf{w}} f(\mathbf{x}; \mathbf{w})$$

- ▶ Where η is the learning rate, we assume it to be 1 for now
- ▶ Here \mathbf{w}^t is the column vector of class k at training step t
- ▶ Consider a class k sample, the change in the parameters will be
- ▶ $\mathbf{w}^{t+1} = \mathbf{w}^t - (\hat{\mathbf{y}}_k - \mathbf{1}) \mathbf{x} = \mathbf{w}^t + (1 - \hat{\mathbf{y}}_k) \mathbf{x}$
- ▶ As $(1 - \hat{\mathbf{y}}_k)$ is always > 0 , the update of \mathbf{w}^t is like adding a scale of \mathbf{x} to it. It is like moving \mathbf{w}^t towards \mathbf{x} .
- ▶ When \mathbf{w}^t is already pointing in the similar direction as \mathbf{x} , this move will increase the magnitude of \mathbf{w}^t . Thus, it is general to observe that, when there are more training samples of class k, the magnitude of column vector of class k is larger.

Update the parameters

- ▶ The parameters will be updated according to

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{d}{d\mathbf{w}} f(\mathbf{x}; \mathbf{w})$$

- ▶ Where η is the learning rate, we assume it to be 1 for now
- ▶ Here \mathbf{w}^t is the column vector of class k at training step t
- ▶ Consider a class i ($i \neq k$) sample, the change in the parameters will be
- ▶ $\mathbf{w}^{t+1} = \mathbf{w}^t - \hat{y}_k \mathbf{x}$
- ▶ As \hat{y}_k is always > 0 , the update of \mathbf{w}^t is like subtracting a scale of \mathbf{x} to it. It is like moving \mathbf{w}^t away from \mathbf{x} .

Update the parameters


- ▶ The parameters will be updated according to

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \eta \frac{d}{d\mathbf{w}} f(\mathbf{x}; \mathbf{w})$$


- ▶ Where η is the learning rate, we assume it to be 1 for now
- ▶ Here \mathbf{w}^t is the column vector of class k at training step t
- ▶ Consider a class i ($i \neq k$) sample, the change in the parameters will be
- ▶ $\mathbf{w}^{t+1} = \mathbf{w}^t - \hat{y}_k \mathbf{x}$
- ▶ As \hat{y}_k is always > 0 , the update of \mathbf{w}^t is like subtracting a scale of \mathbf{x} to it. It is like moving \mathbf{w}^t away from \mathbf{x} .

The XOR problem

- ▶ Can you deduce the values of w and b for the XOR case?




AND function




OR function

The XOR problem


- Consider the very simple task of computing simple logical functions of two inputs,



AND



OR



XOR




The XOR problem

- ▶ XOR is not a linearly separable function, it cannot be solved by logistic regression.
- ▶ Solution: neural network

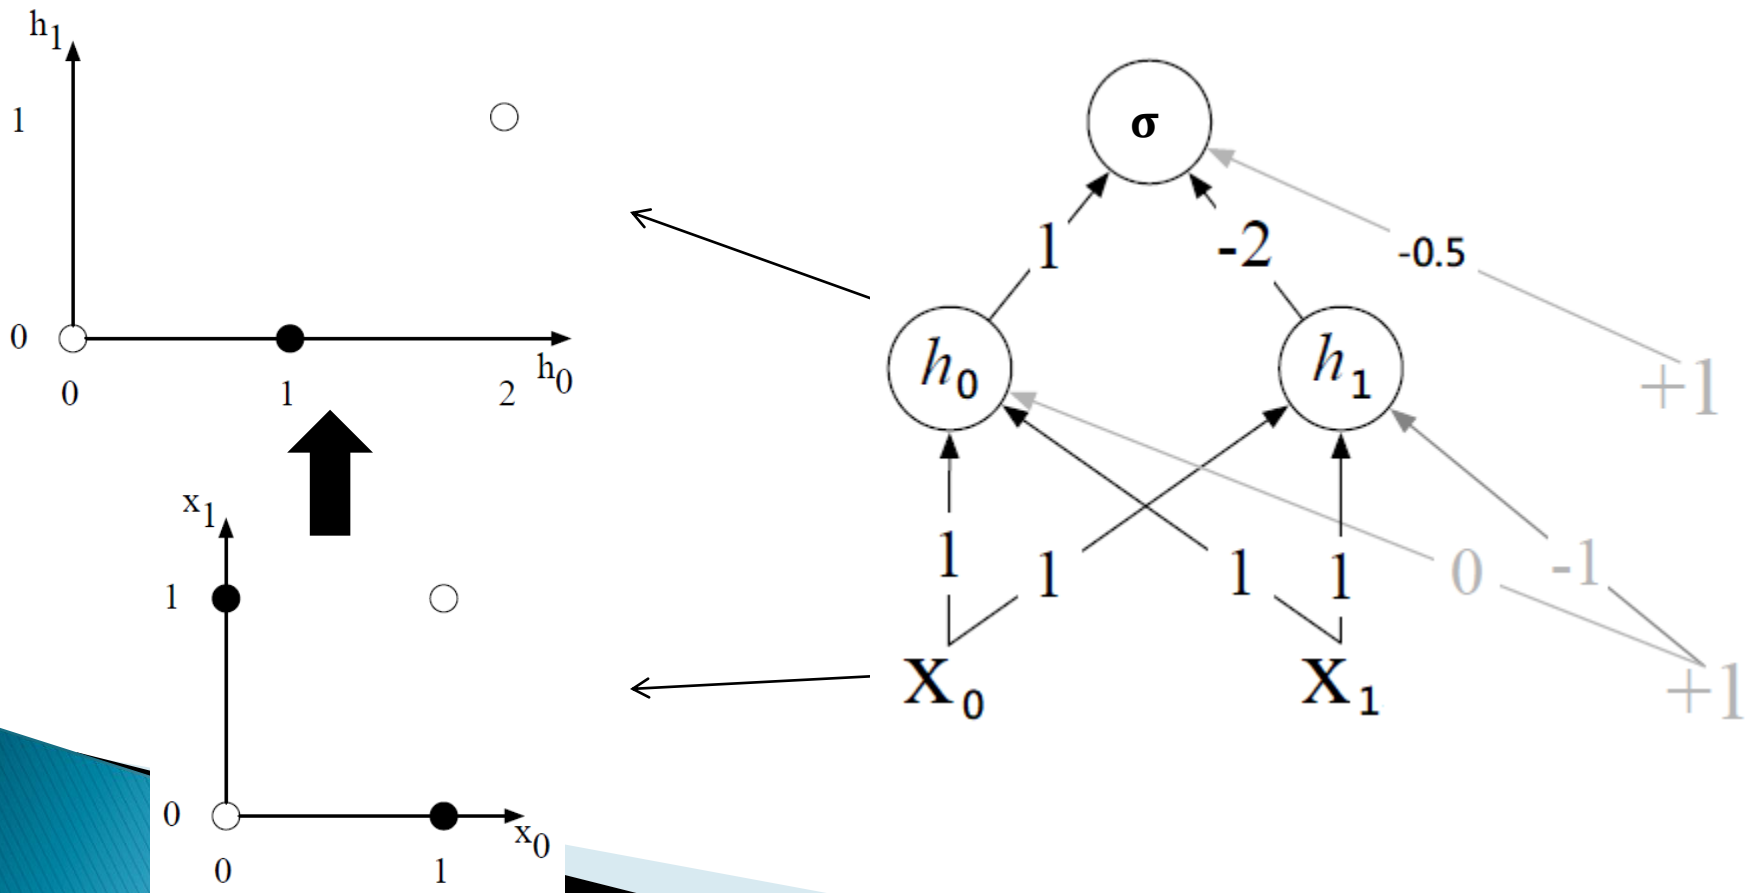
Solution for the XOR problem

- ▶ Using a 2-layer neural network with ReLU-based units




Solution for the XOR problem

- ▶ The hidden layer will learn to form useful representation





An example of neural unit




Activation functions

- ▶ Nowadays, the sigmoid is not commonly used as an activation function.
- ▶ Instead, people use
 - Tanh function
 - Rectified linear unit (ReLU)



Without the non-linear layer

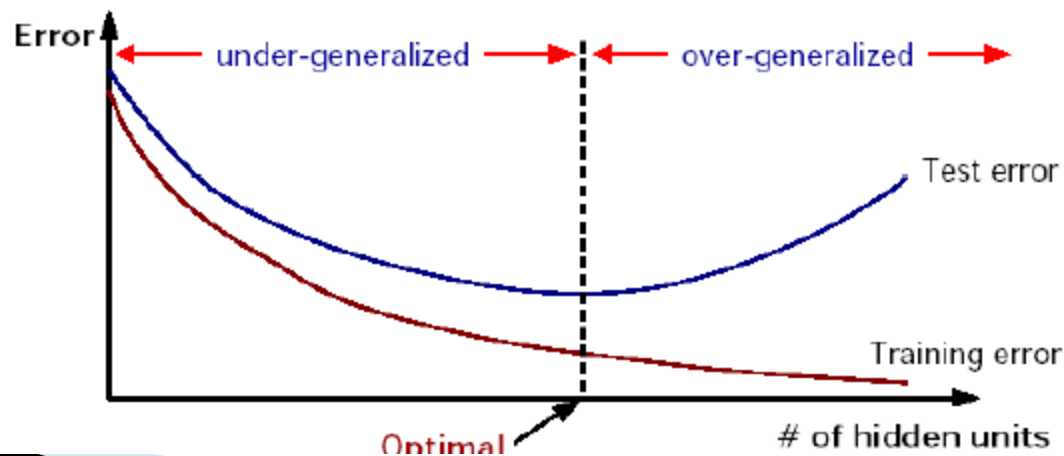
- ▶ A network formed by many layers of purely linear units can always be reduced to a single layer of linear units with appropriate weights.






What is a Deep neural network?


- ▶ Why need networks with > 2 hidden layers?
 - by using extra layers we might find a network with fewer weights in total while still achieving the same level of accuracy
- ▶ How many hidden units?
 - not too many nor too few



Classification with a neural network



Backpropagation




State alignment for training

- Given a 1-sec audio, resulting in 100 MFCCs, with a transcription of word “computer”




Feature Extraction

Observations:




Ground truth
label:

1	1	1	1	1	0	0	0	0	0	0	0	0	0
0	0	0	0	0	1	1	1	1	0	0	0	0	0
0	0	0	0	0	0	0	0	0	1	1	1	1	1



Feed-forward and backpropagate




Soft alignment vs. hard alignment

- Given a 1-sec audio, resulting in 100 MFCCs, with a transcription of word “computer”




Feature Extraction

Observations:



Ground truth
label:

1	1	1	0.9	0.7	0.5	0.2	0.05	0	0	0	0	0	0
0	0	0	0.1	0.3	0.5	0.8	0.9	0.8	0.4	0.2	0	0	0
0	0	0	0	0	0	0	0.05	0.2	0.6	0.8	1	1	1






Architectures of neural network

- ▶ Fully connected neural network
- ▶ Convolutional neural network (CNN)
- ▶ Recurrent neural network (RNN)
 - Long short-term memory (LSTM)

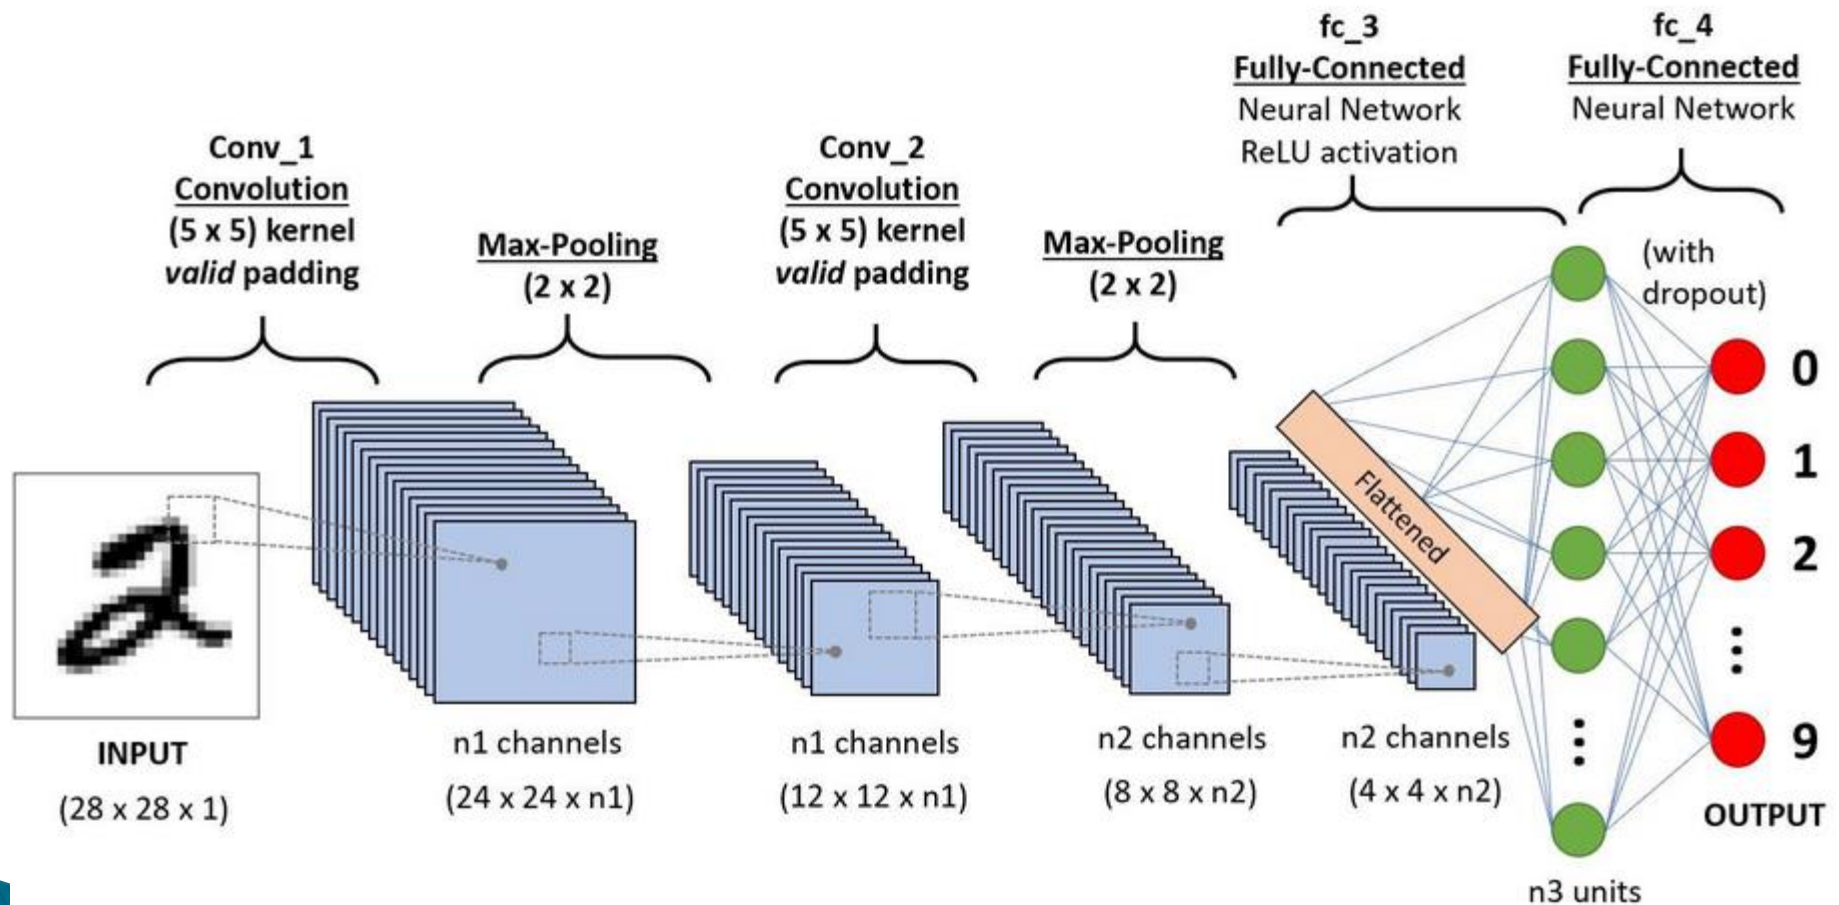
Fully connected neural network



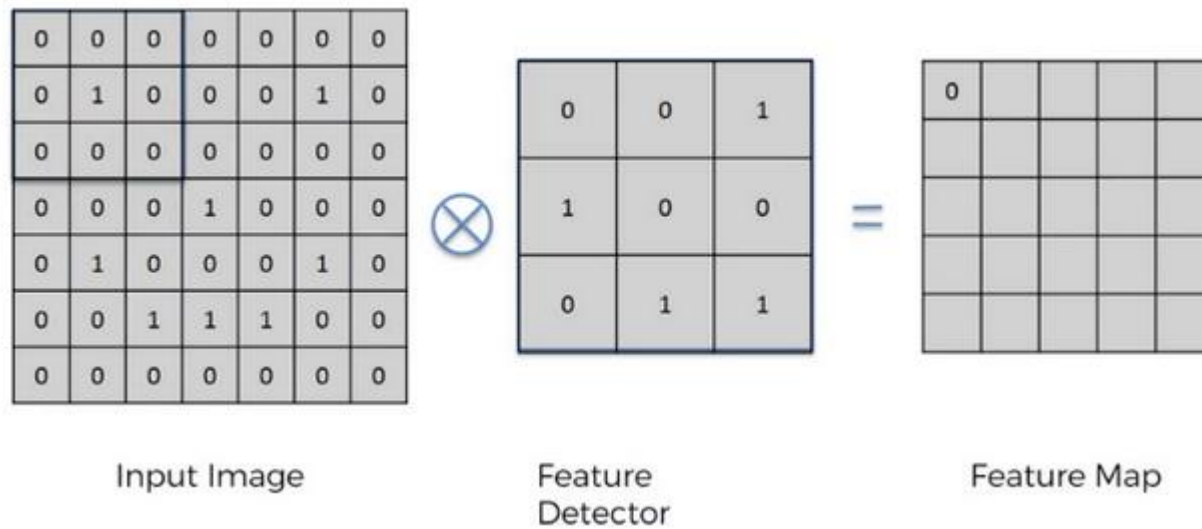
- ▶ How to handle when the input is 2D (e.g. an image)?



Convolutional neural network (CNN)

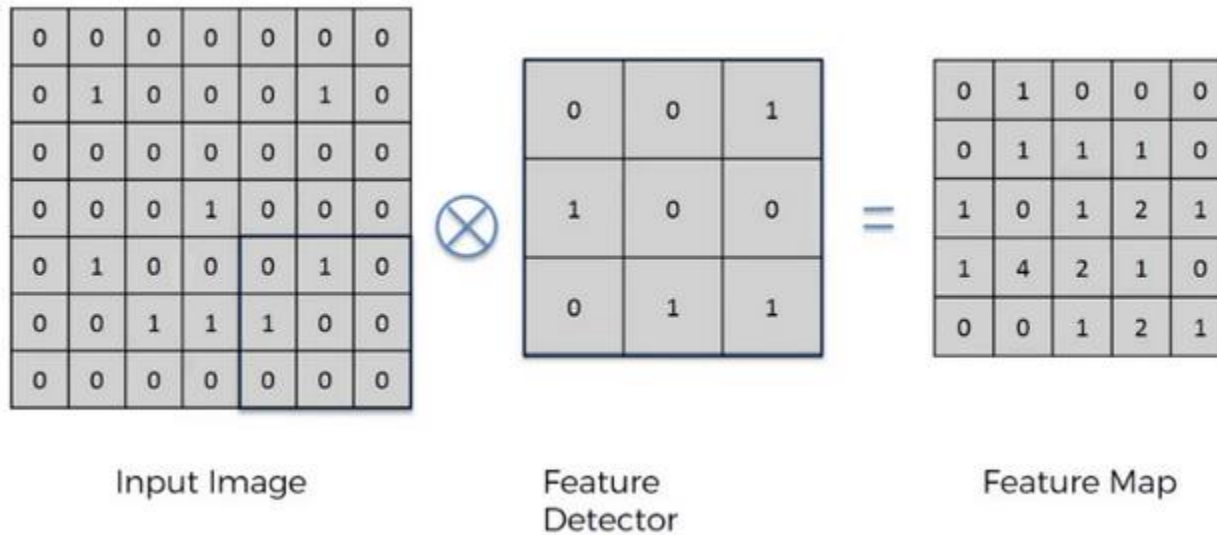


Convolutional layer




- Figures from <https://medium.com/jameslearningnote>

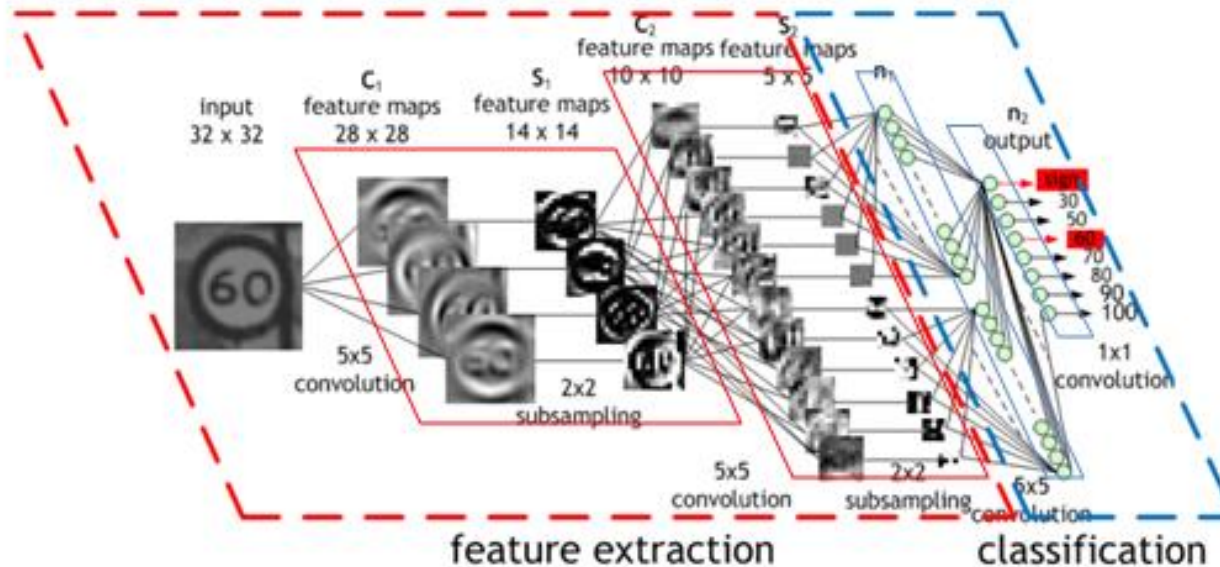
Convolutional layer




Feature maps



CNN as feature extraction




The spectrogram



CNN in speech recognition

- ▶ Treat the spectrogram as an image
- ▶ Convolving along the time axis and the frequency axis





Reading list

- ▶ Speech and Language Processing, version 3
 - Chapter 5 and 7